

CSCI 6461

Computer Systems

Architecture

Addressing Memory

ELF Memory Map

The segments of a program in execution

Executable and Linkable Format (ELF)

- ELF (Executable and Linkable Format) is a file format for executables, libraries, and object code, defining memory layout and linking.
- An ELF (Executable and Linkable Format) number of sections varies depending on the file's purpose (e.g., executable, shared library, or object file) and how it was compiled.
- When a program is loaded into memory it gains additional "dynamic" segments

ARM32 ELF and unix-like systems

- ARM32 binaries are typically compiled into ELF files,
 - this define how the code and data are organized for a executable or linkable file
- ELF is a flexible format supports various processor-specific details,
 - such as ARM's instruction set and memory alignment requirements,
 - machine type field (e.g., EM_ARM for ARM architecture).

C program example


From the course
repo - a simple C
program to show
the various ELF
segments using
objdump -c -S
a.out

```
/*  
 * simple.C  
 * a file to show the layout of objects  
 * in the a.out executable  
 * using objdump -S -C  
 * */  
#include <stdio.h>  
  
// initialized global variable (r/w)  
int global_var1 = 100;  
  
// uninitialized global variable (r/w)  
int global_var2;  
  
// initialized global variable (r/w) string  
char global_var3[] = "This is a global string";  
  
int main(int argc, char* argv[]) {  
    int local_var = 99;  
    // the reference to printf should be mentioned in the .dynstr  
    // the string itself ("Globals are ...") is r/o in .data  
    printf("Globals are %d %d %s\n", global_var1, global_var2, global_var3);  
    return 0;  
}
```

The Linux Virtual Memory Map

6

When the ELF executable becomes a **process**, it gains a number of additional memory resources, such as stack, heap and share library region



ARM32 Memory Map Sections

- **Kernel Space** : 1 GB (1024 MB) reserved for the kernel.
- **Stack**: Thread stacks, grows downward
- **Memory Mapping** : Shared libraries, mapped files.
- **Heap**: Dynamic memory, grows upward
- **BSS**: Uninitialized global/static variables (zeroed).
- **Data**: Initialized global/static variables.
- **Text** (Code): Program's executable code (read-only).

Text and Data Segments

8

- The **text** segment stores the machine language program.
 - In addition to code, also literals (constants) and read-only data
- The global **data** segment stores global variables that, in contrast to local variables, can be accessed by all functions in a program (Read/Write segment)

Dynamic Data Segment

9

- The **dynamic data** segment holds the **stack** and the **heap**.
 - The data in this segment is not known at start-up but is dynamically allocated and deallocated throughout the execution of the program
- On start-up, the operating system sets up the stack pointer (sp) to point to the top of the stack.
 - The stack typically grows downward.
- The stack includes temporary storage and local variables, such as arrays, that do not fit in the registers

Addressing Memory

Modes for accessing memory

What are 'Addressing Modes'?

11

- Data is rarely hard-coded into an instruction using immediates (eg `mov r0, #2`).
- We nearly always read our data from physical memory (RAM),
 - then we can use the data, and
 - write it back to physical memory
- There are different methods (or modes) we can use to achieve the correct outcome

Memory Addressing Modes

12

- Register indirect
 - `ldr r1, [r0]`
- Register indirect with offset
 - `ldr r1, [r0, #4]`
- Double Register Indirect
 - `ldr r1, [r0, r2]`

in these examples, `r0` is the **base** register

Initializing the base pointer in C

In C, we are all familiar with the following approach:

So `ptr` is our base pointer (or base address)

```
#include <stdio.h>

int array[6] = { 2,3,11,4,55,6 };

int main(int argc, char* argv[]) {
    int *ptr = array;
    *(ptr += 4) += 10;
    printf("New number: %d\n", *ptr);
    return 0;
}
```

<https://github.com/jzburns/csci-comp-arch/blob/master/C/simple.C>

Initializing the base register in ARM

```
//Step 1 - declare a .data section (like in the ELF)
//Step 2 - add your integer array
//Step 3 - pick a base register (eg, r0) and use the =
//operator to load the address
```

```
.data
    array: .word 2,3,11,4,55,6 int array[6] = { 2,3,11,4,55,6 };
ldr r0, =array int *ptr = array;
```

Register Indirect Addressing

15


ARM has register indirect addressing e.g. loading a register from a memory location:

```
ldr r2, [r0]
```

We can think of r_0 as being the address of the **first** element in the array pointed to by r_0

Register Indirect Addressing

16



- It takes only a few bits to select a register (4 bits in the case of ARM ... r₀-r₁₅)
- A register can (typically) hold an arbitrary address (32 bits in the case of ARM)

Indirect with Offset

17

ARM allows offsets of **12 bits** in ldr/str

This may be added or subtracted, for example:

```
ldr r0, [ r1, #8 ]
```

```
str r3, [ r6, #-64 ]
```


This provides a range of $\pm \sim 4$ Kbytes around a 'base' register - adequate for most purposes.

Indirect with Offset

18

The address is calculated from a register value and a literal

The register specifier is just a few bits
The offset can be 'fairly small'
With one register 'pointer' any of several variables in nearby addresses may be addressed



Address Arithmetic

19

We can operate on registers, so we can:

- store/load/move addresses
- do arithmetic to calculate addresses

Rather than using extra add instructions, we often use

Base + Offset Addressing

- address addition done within the operand


Double Register Indirect

20

ARM also allows offsets using a second register
This may be added or subtracted, for example:

```
ldr r0, [ r1, r2 ]
```

```
str r3, [ r6, -r4 ]
```



PC + Offset Addressing

21

start:

```
ldr r0, =b  
ldr r1, [ r0 ]  
add r2, r1, #5  
str r2, [ r0 ]
```

The PC is the address of the current instruction plus 8 bytes (2 words). This means

ldr r0, =b is actually: ldr r0, [pc, #8]

What effect does the code have?

.data

```
b: .word 20  
c: .word 30
```

[Why does the ARM PC register point to the instruction after the next one to be executed? - Stack Overflow](#)

Pointer arithmetic

```
*(ptr += 4) += 10;
```

How can we implement this in ARM32 assuming r0 is our base register? Let's try:

```
add r0, r0, #16
```

```
ldr r5, [ r0 ]
```

```
add r5, r5, #10
```

```
str r5, [ r0 ]
```

Iteration over an array

Now we can use `lsl`

Adding 10 to each element in an array

```
.global _start
_start:
    ldr r0, =array1      // array base pointer
    mov r1, #0            // offset
    mov r2, #0            // loop counter

loop:
    cmp r2, #6
    bge end
    ldr r3, [r0, r1]      // read base + offset to reg
    add r3, r3, #10        // add 10 to the value
    str r3, [r0, r1]      // write reg to base + offset
    add r2, r2, #1          // inc loop
    lsl r1, r2, #2         // calc offset
    b loop
end:
    svc #2

.data
array1: .word 2,3,11,4,55,6
```

A note on Incrementing & offsets

Effect of LSL R2, R1, #2

Itr	R1	R2	Dec
0	0000	0000	0
1	0001	0100	4
2	0010	1000	8
3	0011	1100	12
4	0100	10000	16

Operations within an instruction

Iteration optimization from the previous example:

```
ldr r3, [ r0, r2, lsl #2 ]  
add r2, r2, #1
```

Thus we can eliminate the `lsl r1, r2, #2`

The highlighted code does **not** change the value of `r2`

Register Byte Instructions

ldrb / strb

Strings

28

- Java:
 - `String message= "Hello";`
- ARM:
 - `.data`
 - `message: .asciz "Hello World!"`

Iterating over an array of "chars"

29

Suppose we want to iterate over each character individually:

```
for (int i= 0; i<message.length(); i++) {  
    // do something with this:  
    message.charAt(i);  
}
```



Accessing characters: LDRB, STRB

30

```
ldr r1, =message  
ldrb r0, [r1]      // fetch 1st byte  
ldrb r0, [r1,#1]   // fetch 2nd byte  
ldrb r0, [r1,#2]   // fetch 3rd byte
```

get next character from the string

31

```
for (int i= 0; . . . ; i++) {  
    System.out.print(message.charAt(i));  
}
```

Use a second register (e.g. R2) to hold i around loop:

First attempt, no loop

32


```
mov r2, #0           //int i= 0
ldr r0, =message    //beginning address of 'message'
...
ldrb r0, [r1,r2]    // message.charAt(i) = r1+r2
add r2, r2, #1      // i++
...
...
```

ldrb Load Register Byte (immediate) loads a byte from memory, zero-extends it, and writes the result to a register.

get next character from the string

33

Part way
through
execution



get next character from the string

34

optimisation: avoid using both r1 and r2 for addresses
actually change the address in r1

```
ldr r0, =message
```

...

```
ldrb r1, [ r0 ]
```

```
add r0, r0, #1
```

Indexing Options

3 main types: offset, pre and post

Indexing modes

In addition to scaling the index register, ARM provides

- offset addressing
- pre-indexed addressing
- post-indexed addressing

Mode	ARM Assembly	Address	Base Register
Offset	LDR R0, [R1, R2]	R1 + R2	Unchanged
Pre-index	LDR R0, [R1, R2]!	R1 + R2	R1 = R1 + R2
Post-index	LDR R0, [R1], R2	R1	R1 = R1 + R2

get next character from the string

37

- Optimisation: avoid using `add r1, r1, #1`
- Change the address in `r1` using “**post-indexed**” operand form

```
ldr r1, =message
```

```
ldrb r0, [r1],#1
```

while not at end of string

38

Every Java/C String knows how long it is. But how do we know in assembly when the end of the string is reached?

```
.data
```

```
message: .asciz "Hello World!"
```

Then iterate over the array of characters until 0 is reached.
In C/C++ this is sometimes referred to as the **null** character.

Whole loop

39

```
ldr r0, =message
loop:
    ldrb r1, [ r0 ],#1
    cmp r1, #0
    beq end
    b loop
end: // end here
.data
    message: .asciz "Hello World!"
```

Further Examples

40

(simple) indirect:

```
ldr r0, [ r1 ] ; r0 ← value at [r1], r1 unchanged
```

(base +) offset:

```
ldr r0, [ r1, #4 ] //r0 ← value at [r1+4], r1 unchanged
```

```
ldr r0, [ r1, r2 ] //r0 ← value at [r1+r2], r1 unchanged
```

post-indexed:

```
ldr r0, [ r1 ], #4 //r0 ← value at [r1], r1 ← r1+4
```

```
ldr r0, [ r1 ], r2 // r0 ← value at [r1], r1 ← r1+r2
```

pre-indexed:

```
ldr r0, [ r1, #4 ]! // r0 ← value at [r1+4], r1 ← r1+4
```

```
ldr r0, [ r1, r2 ]! //r0 ← value at [r1+r2], r1 ← r1+r2
```