

# CSCI 6461 Computer Systems Architecture

Pipelined processor

# Processing Stages

We design a pipelined processor by subdividing the single-cycle processor into five pipeline stages:

- Fetch
- Decode
- Execute
- Memory
- Writeback

# Simultaneous Execution

- They are similar to the discrete stages in the single-cycle processor used to perform `ldr`
- Five instructions can execute simultaneously,
  - one in each stage.
- As we will see, this introduces significant performance efficiencies
  - but introduces the idea of a **hazard**

# The 5 Stages

- **Fetch** : The processor reads the instruction from instruction memory.
- **Decode** : The processor reads the source operands from the register file and decodes the instruction to produce the control signals.
- **Execute** : The processor performs a computation with the ALU.
- **Memory** : The processor reads or writes data memory.
- **Writeback** : The processor writes the result to the register file, when applicable.

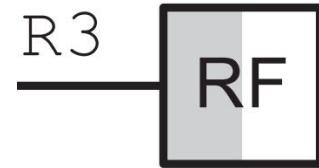
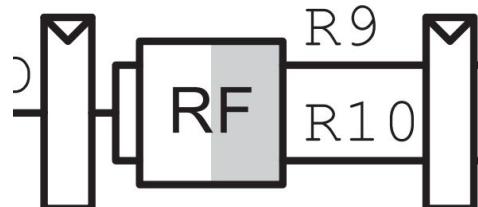
# Simple Program

Our test case for pipeline analysis:

```
ldr r2, [ r0, #40 ]
add r3, r9, r10
sub r4, r1, r5
and r5, r12, r13
str r6, [ r1, #20 ]
orr r7, r11, #42
```

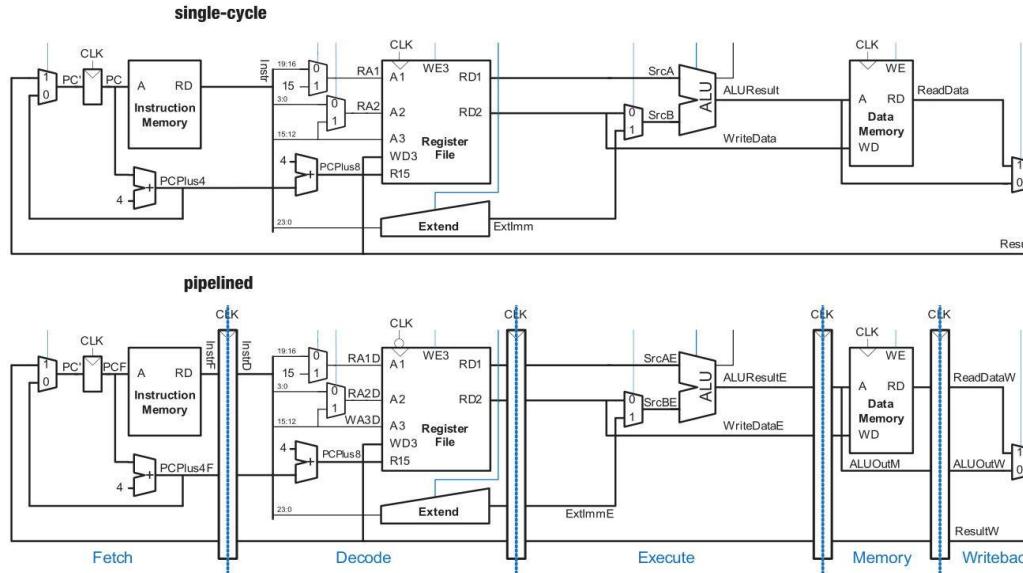
# A note on shading

- Some resources, such as the register file (RF), can support read & write in the same phase.
  - shading of the RF on the right side means it is **read** in the phase
  - and shading on the left side means it is **written** in the phase
- No shading means idle



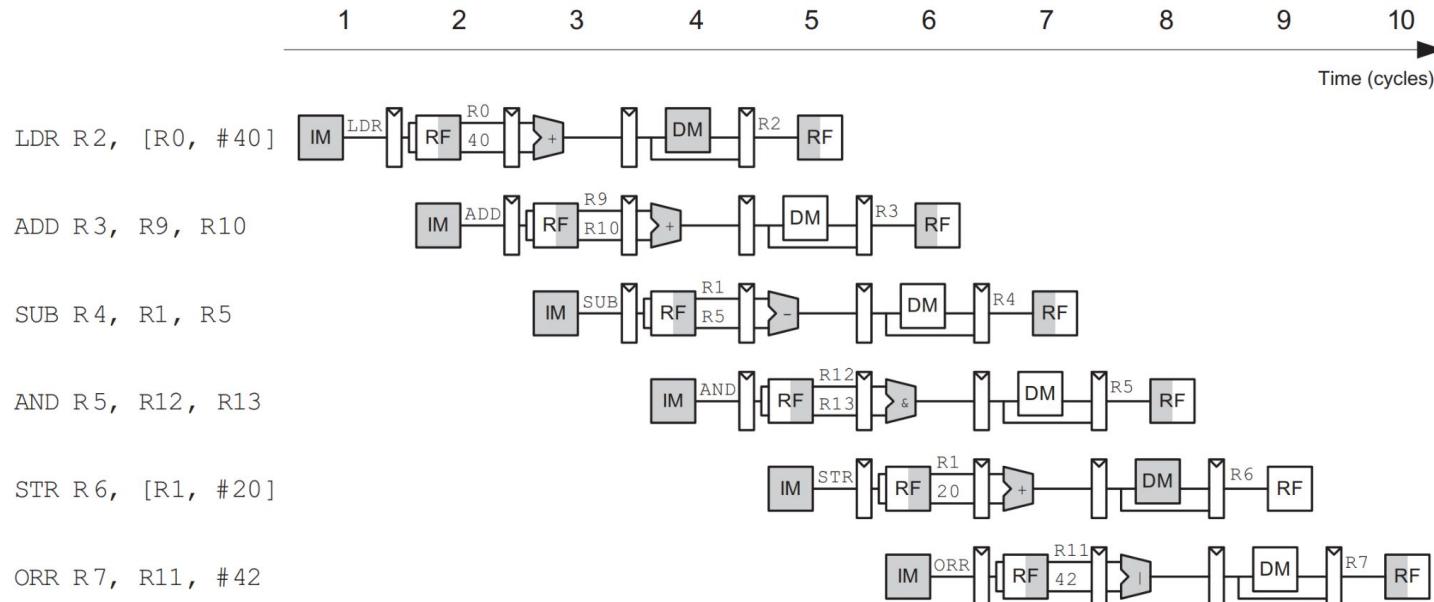
# Chop the single-cycle into 5 stages

The pipelined datapath: chop the single-cycle datapath into five stages and separate them by pipeline registers.



Signals are given a suffix (F, D, E, M, W) to indicate the stage

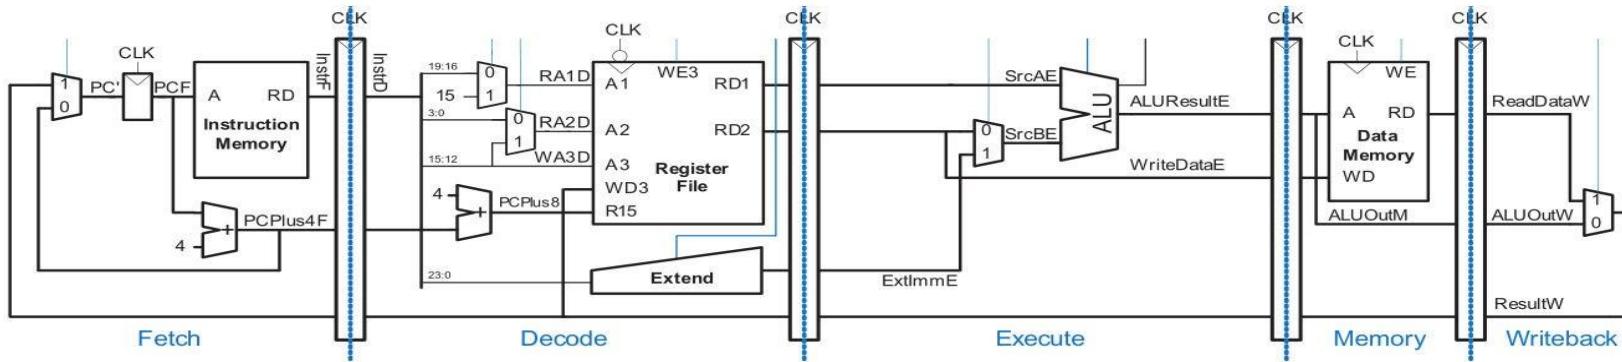
# Theoretical Instruction Pipeline



- In the pipelined processor, the register file is written in the 1st part of a cycle and read in the 2nd part.
- This way, data can be written and read back within a single cycle.

# Register File Write Stage

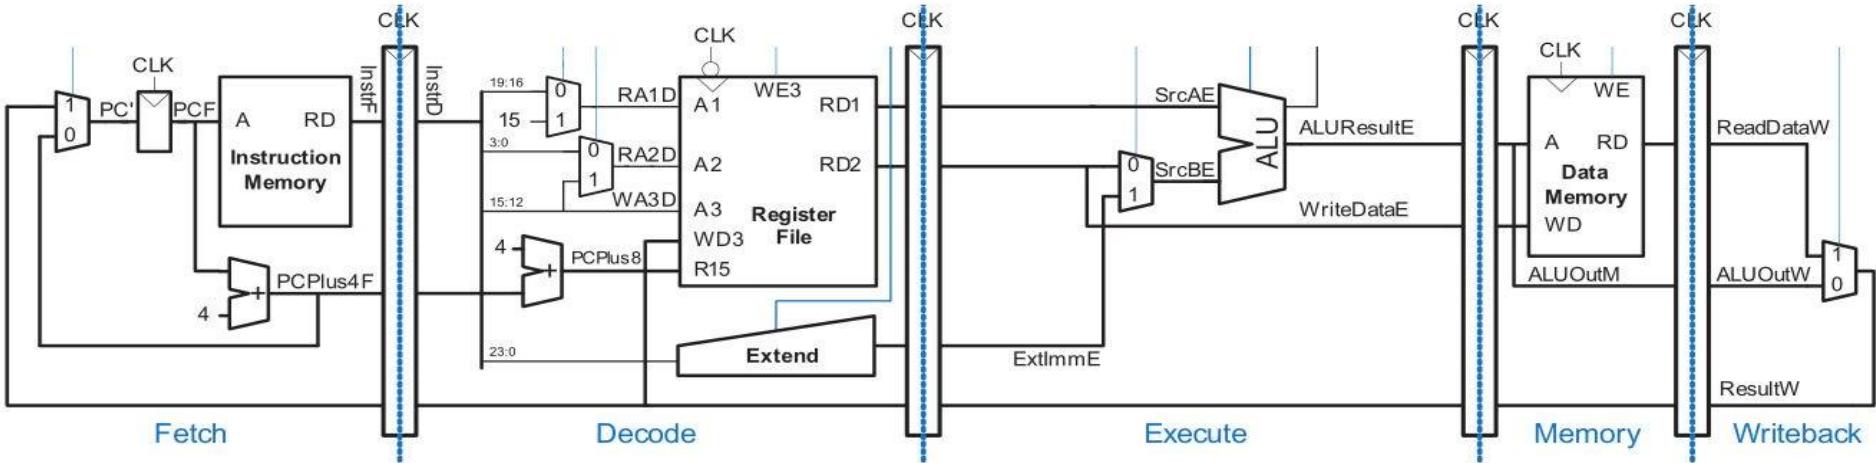
- The register file is read in the **Decode** stage and written in the **Writeback** stage.
- It is placed in the **Decode** stage, but the write data comes from the **Writeback** stage.



# CLK cycle write and reads

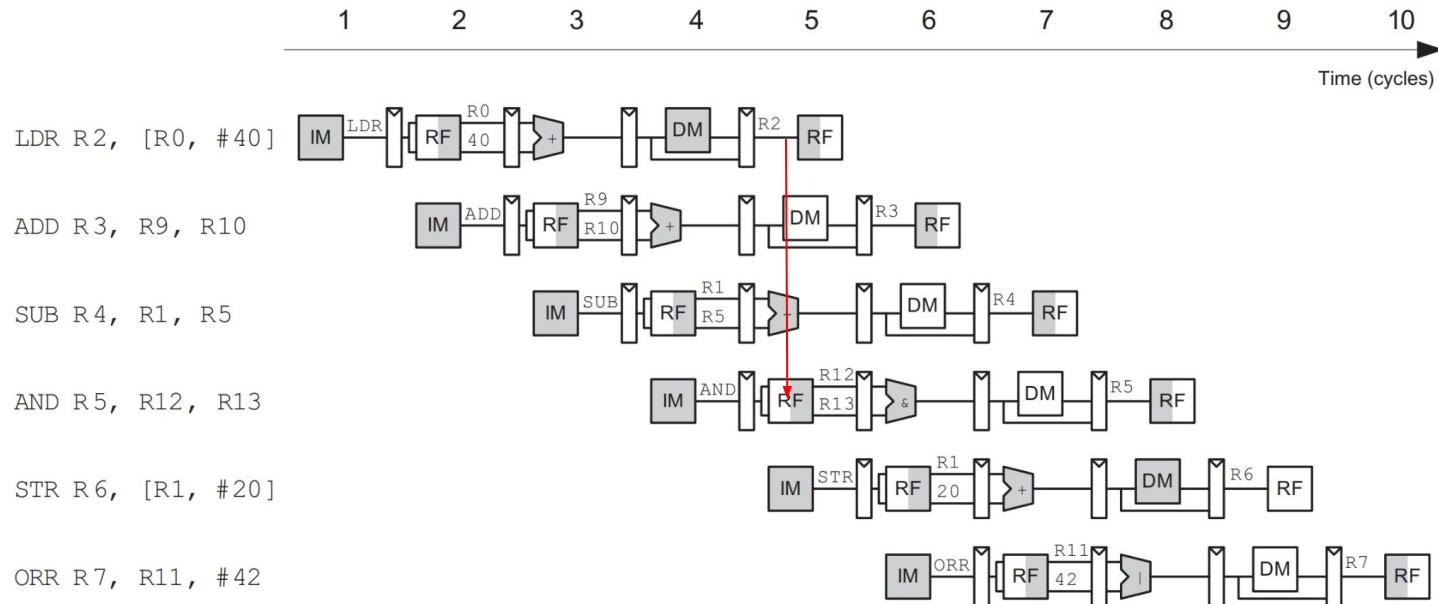
- The register file in the pipelined processor writes on the falling edge of CLK
  - so that it can write a result in the first half of a cycle and read that result in the second half of the cycle
  - for use in a subsequent instruction.
- A critical issue in pipelining is that all signals associated with a particular instruction must advance through the pipeline in unison

# Datapath Stage Errors



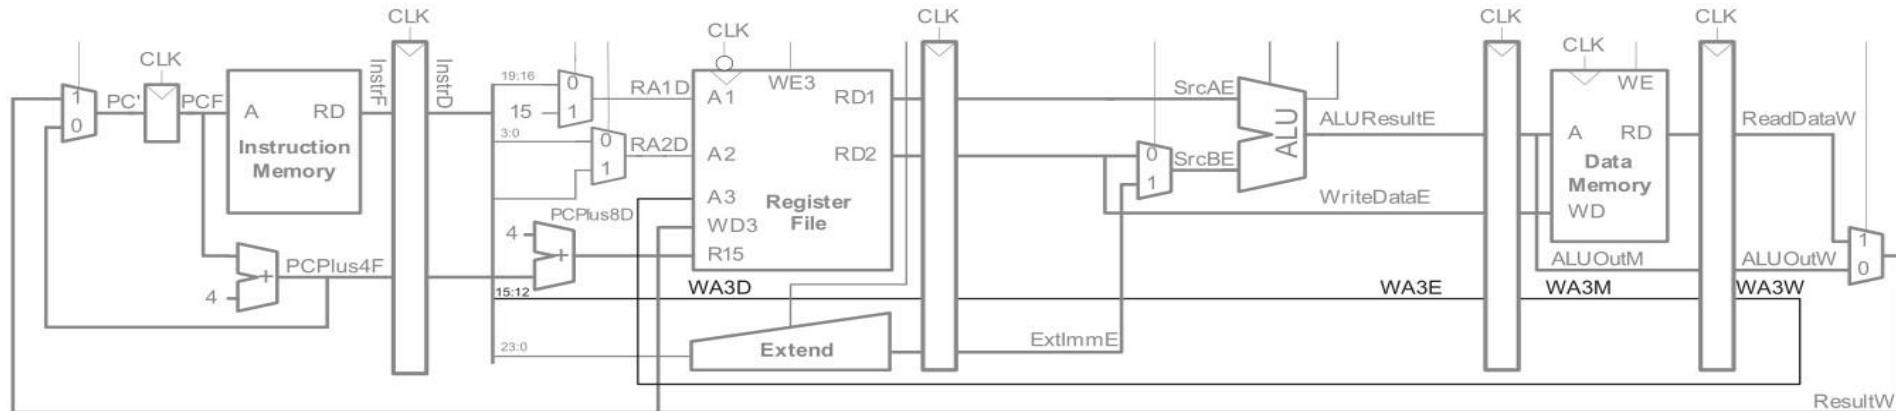
- Register file write logic operate in the **Writeback** stage.
- The data value comes from **ResultW**, a **Writeback** stage signal.
- But the write address comes from  $\text{InstrD}_{15:12}$  (WA3D), which is a **Decode** stage signal.

# Datapath Stage Errors



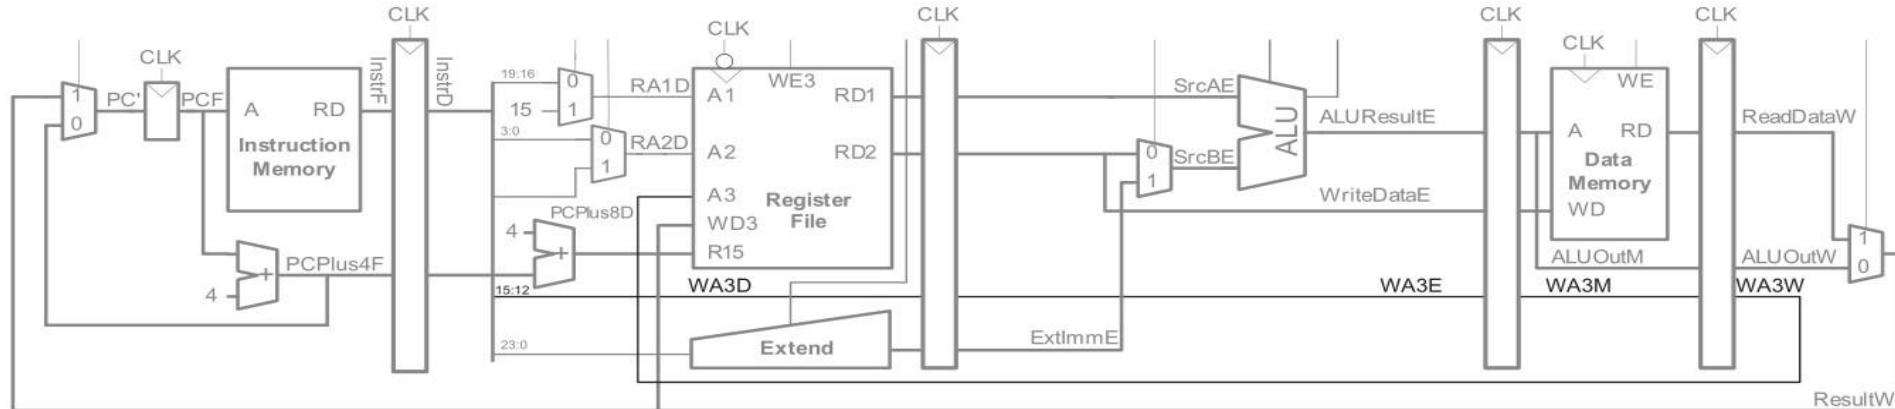
During cycle 5, the result of the `ldr` instruction would be incorrectly written to `r5` (and `r5, r12, r13`) rather than `r2`.

# Corrected WA3D pipelined datapath



- The **WA3D** signal is now pipelined along through the Execution, Memory, and Writeback stages (**WA3E**, **WA3M**, **WA3W**), so it remains in sync with the rest of the instruction.
  - **WA3W** and **ResultW** are fed back together to the register file in the Writeback stage.

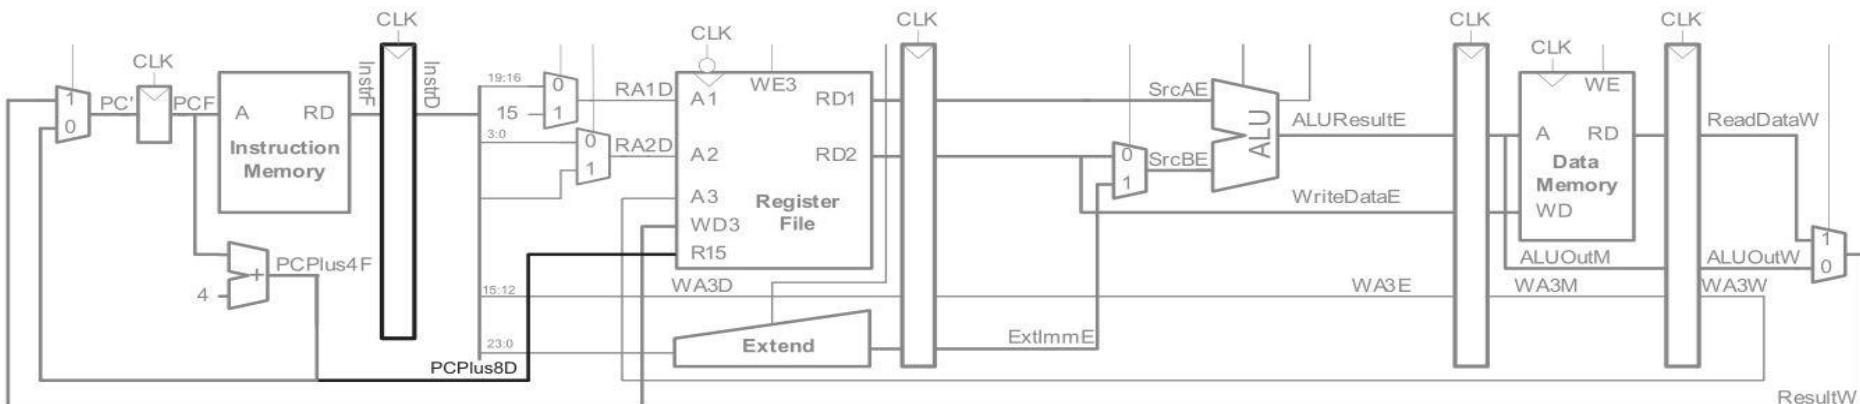
# Incrementing the PC



- Each time the program counter is incremented, **PCPlus4F** is simultaneously written to the PC and the pipeline register between the F and D stages.
- On the subsequent cycle, the value in both of these registers is incremented by 4 again.
- Thus, **PCPlus4F** for the instruction in the Fetch stage is logically equivalent to **PCPlus8D** for the instruction in the Decode stage.

# Optimizing PC incrementation

- Sending this signal ahead saves the pipeline register and second adder.
- Optimized PC logic eliminating a register and adder:



# PC control hazard

PC' logic is also problematic:

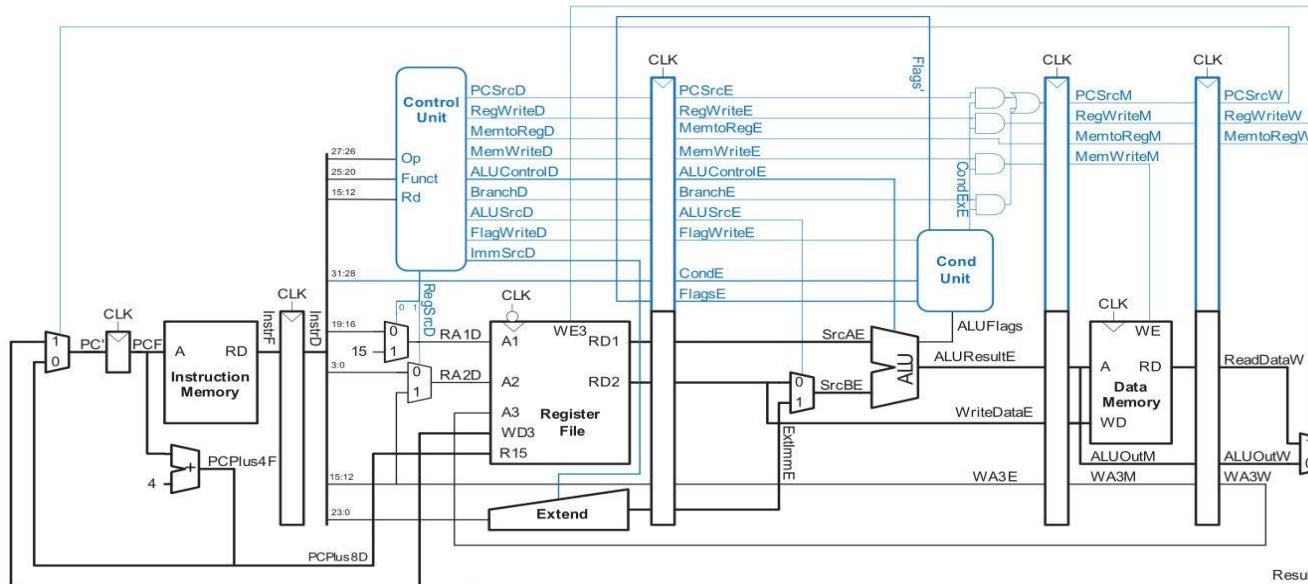
- it might be updated with a Fetch
  - **PCPlus4F** (eg, sequence/branch)
  - **ResultW** (eg, pop {pc})
- This control hazard will be fixed in the next section

# Pipelined Control

# Pipelined Control

- The pipelined processor takes the same control signals as the single-cycle processor and therefore uses the same control unit.
- The control unit examines the **Op** and **Funct** fields of the instruction in the Decode stage to produce the control signals.
- These control signals must be pipelined along with the data so that they remain synchronized with the instruction. The control unit also examines the **Rd** for writes to **r15** (pc).

# Pipelined processor with control



RegWrite must be pipelined into the W stage before it feeds back to the register file, just as WA3 is pipelined.

# Hazards

Consequences of Pipelining in 5 Stages

# Hazard Types

- A **data** hazard occurs when an instruction tries to read a register that has not yet been written back by a previous instruction.
- A **control** hazard occurs when the decision of what instruction to fetch next has not been made by the time the fetch takes place.
- Let's look at a data hazard first

# Data Hazards

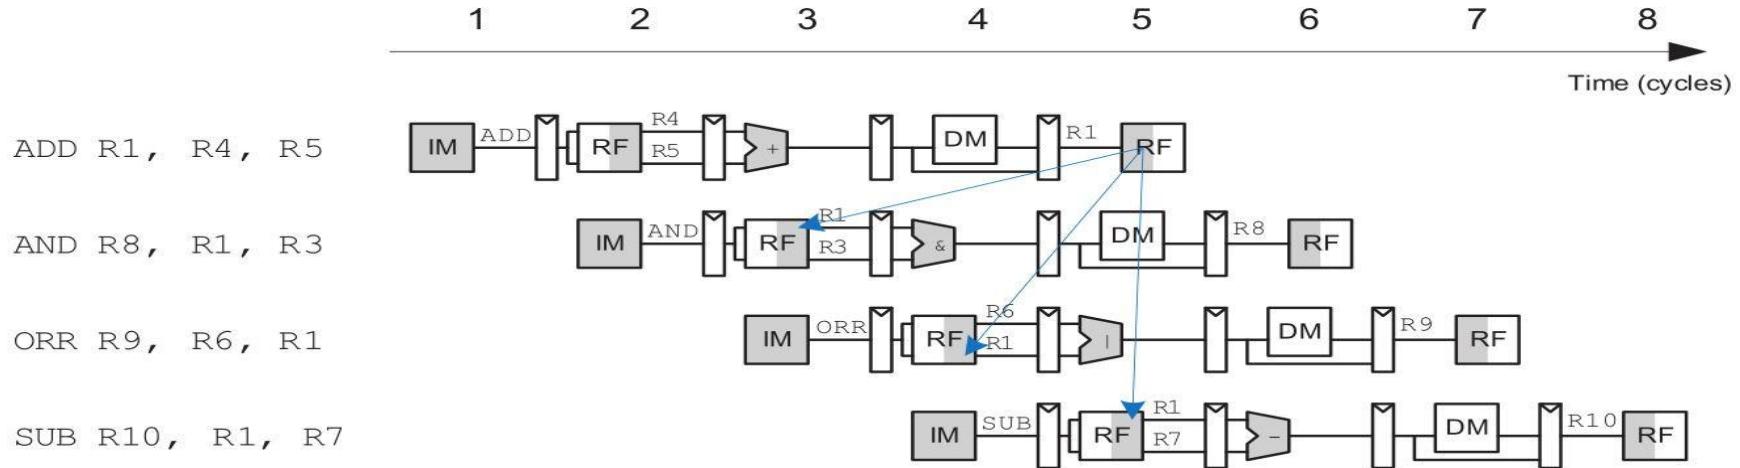
## RAW and Stalls

# Data Hazards

There are several kinds of data hazards to be aware of:

- **Read after Write** (RAW) occurs when the value produced by an instruction is required by a subsequent instruction
- **Write after Write** (WAW) occur when the output register of an instruction is used for write after written by previous instruction

# RAW Hazard



- The add instruction writes a result into R1 in the first half of cycle 5.
- The and instruction reads R1 on cycle 3, obtaining the wrong value.
- The orr instruction reads R1 in the second half of cycle 4, again, obtaining the wrong result
- What about the sub instruction?

# Result Forwarding

- The sum from `add` is computed by the ALU in cycle 3
- We can **forward** the result from one instruction to the next to resolve the RAW hazard
  - **without** waiting for the result to appear in the register file.
- In other situations we may have to **stall** the pipeline to give time for a result to be produced before the subsequent instruction uses the result.

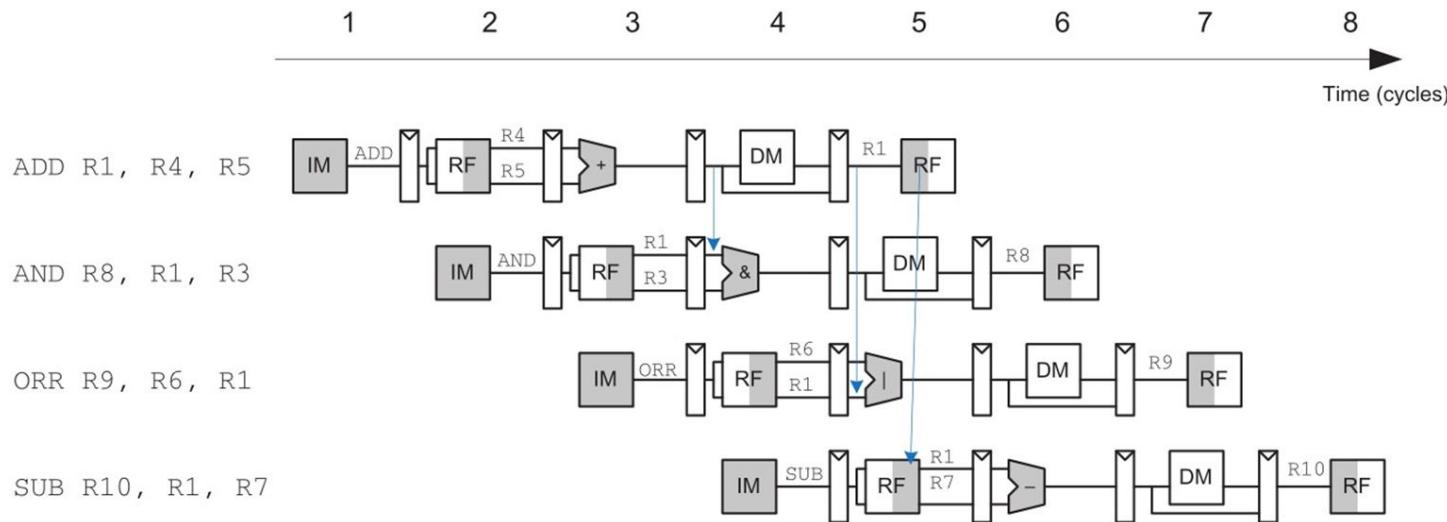
# Hazard Unit

Stage forwarding & multiplexers

# Hazard Units

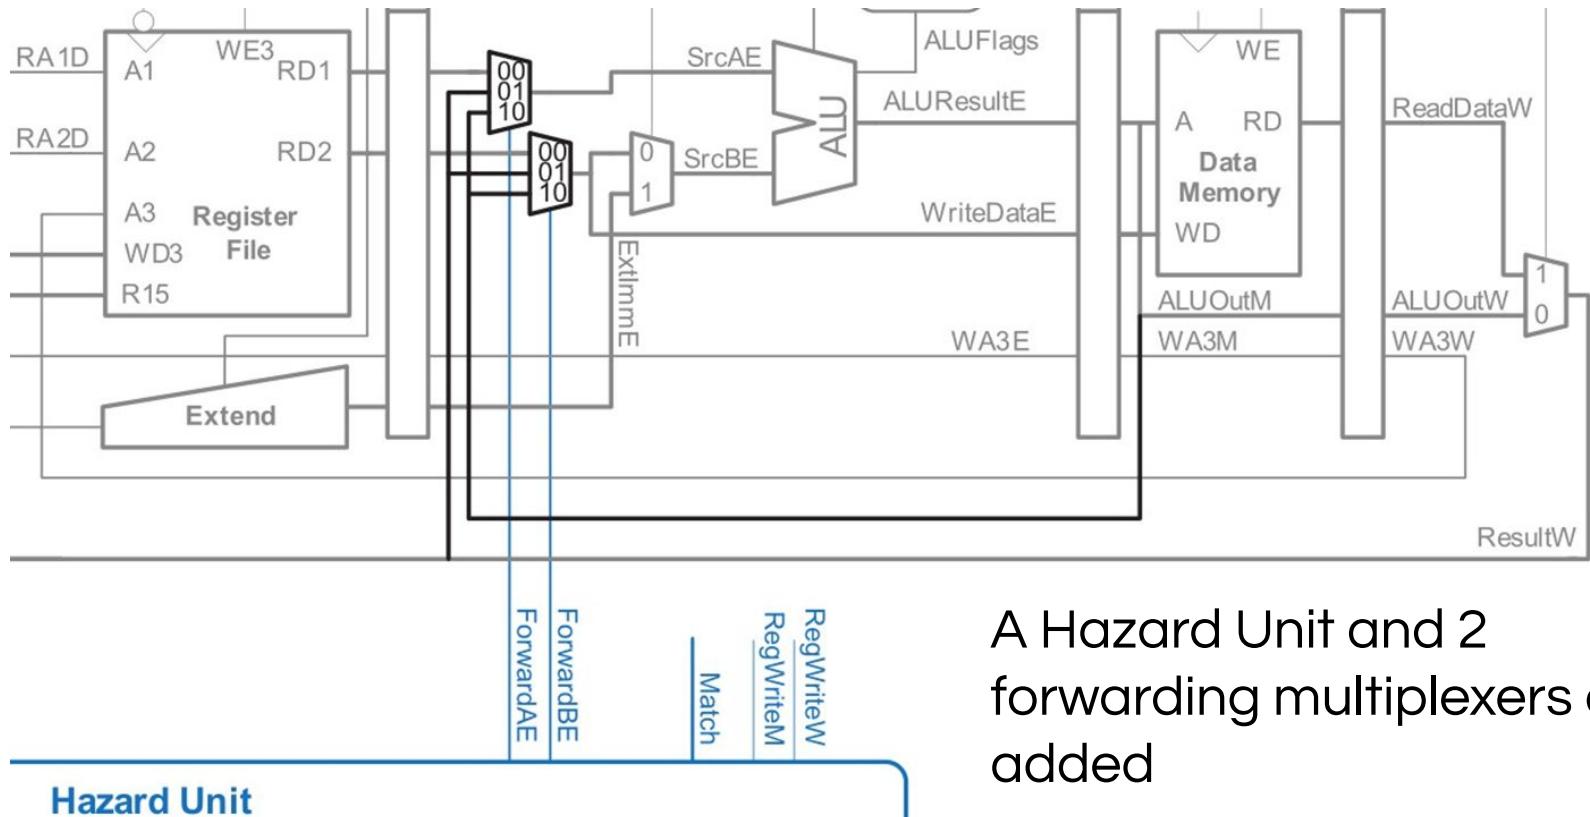
- We will enhance the pipelined processor with a **Hazard Unit**
- The HU detects hazards and handles them appropriately
  - so that the processor executes the program correctly
- The HU must correctly compute the signals to stage forward results

# Add stage forwarding



This requires adding multiplexers in front of the ALU to select the operand from either the register file or the Memory or Writeback stage.

# Architecture with Hazard Unit



# Hazard Unit

- The Hazard Unit receives 4 match signals from the datapath (abbreviated to **Match**)
- that indicate whether the source registers in the Execute stage match the destination registers in the Memory and Execute stages:
  - For example:  $\text{Match\_1E\_M} = (\text{RA1E} == \text{WA3M})$  means a read of A1 in the execute phase matches a write of A3 in the memory phase instruction ahead

```
Match_1E_M = (RA1E == WA3M)
Match_1E_W = (RA1E == WA3W)
Match_2E_M = (RA2E == WA3M)
Match_2E_W = (RA2E == WA3W)
```

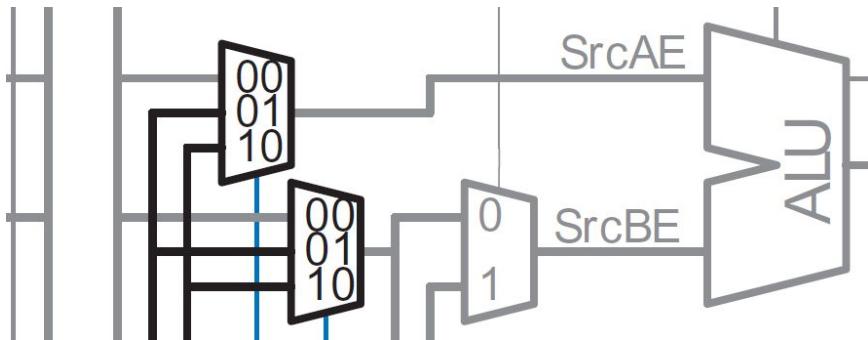
# Hazard Unit

- The Hazard Unit computes control signals for the forwarding multiplexers to choose operands from the register file or from the results in the Memory or Writeback stage (**ALUOutM** or **ResultW** ).
- It should forward from a stage if that stage will write a destination register and the destination register matches the source register.
- If both the Memory and Writeback stages contain matching destination registers, then the Memory stage should have priority, because it contains the more recently executed instruction.

# Forwarding Logic

```
if (Match_1E_M && RegWriteM) // (RA1E == WA3M)
    ForwardAE = 10; // SrcAE = ALUOutM
else if (Match_1E_W && RegWriteW)
    ForwardAE = 01; // SrcAE = ResultW
else ForwardAE = 00; // SrcAE from regfile
```

The forwarding logic for **SrcBE (ForwardBE )** is identical except that it checks **Match\_2E** .



# Match\_1E\_W = (RA1E == WA3W)

- Consider: Match\_1E\_W = (RA1E == WA3W)
- This hazard arises when we have a **RA1 (19:16)** in execute that matches **WA3** in the write phase.
  - add r5, r4, r2 // WA3W
  - sub r7, r8, r9 // Other instr
  - add r6, r5, r1 // RA2E

# Match\_2E\_M = (RA2E == WA3M)

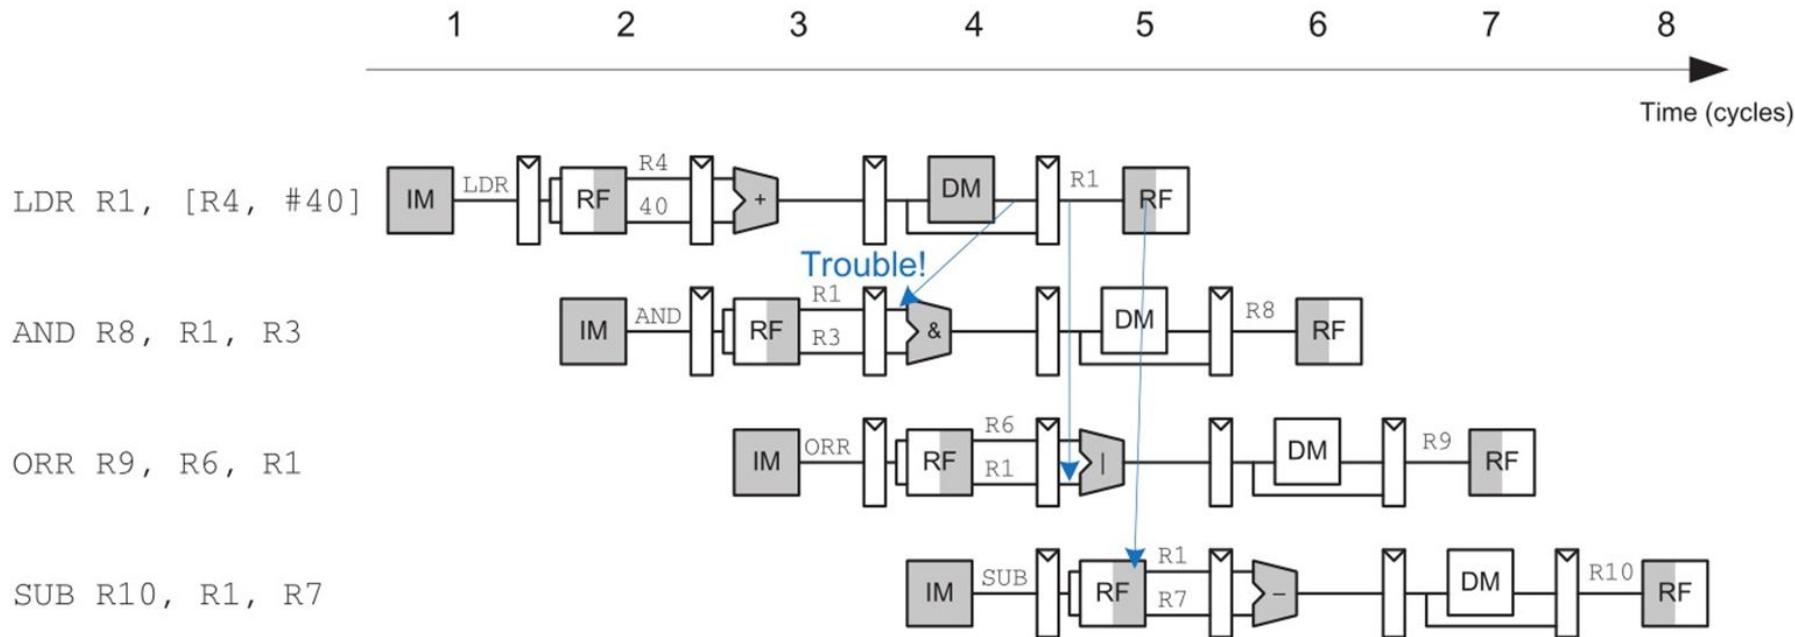
- Consider: Match\_2E\_M = (RA2E == WA3M).
- This hazard arises when we have a **RA2 (3:0)** in execute that matches **WA3** in the memory phase.
  - add r5, r4, r2 // WA3M
  - add r6, r4, r5 // RA2E

This is the tightest RAW dependency the ARM pipeline can still resolve without stalling :  
“one-instruction-apart” case

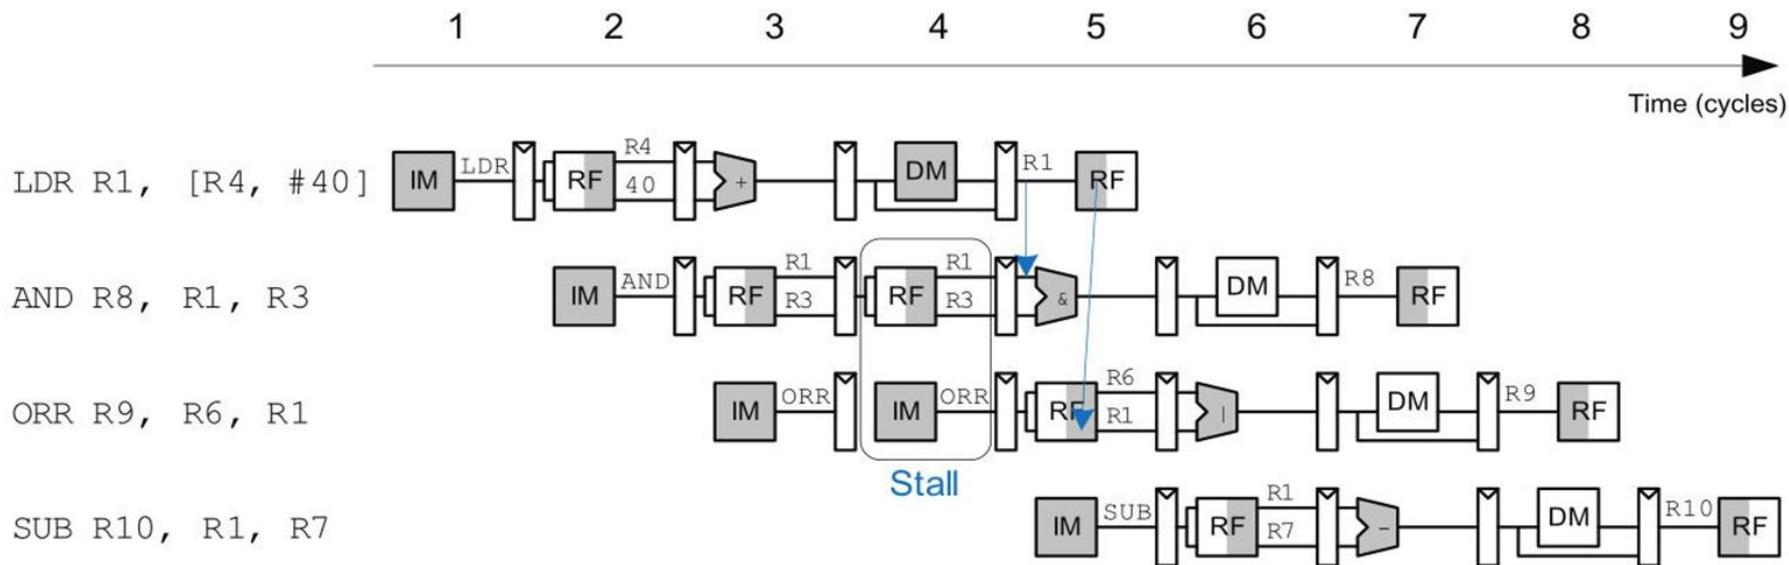
# Trouble forwarding from ldr

- The ldr instruction does not finish reading data until the end of the Memory stage,
  - so its result cannot be forwarded to the Execute stage of the next instruction.
- ldr has a two-cycle latency,
  - because a dependent instruction cannot use its result until two cycles later.
- We cannot solve this hazard with forwarding.

# For example

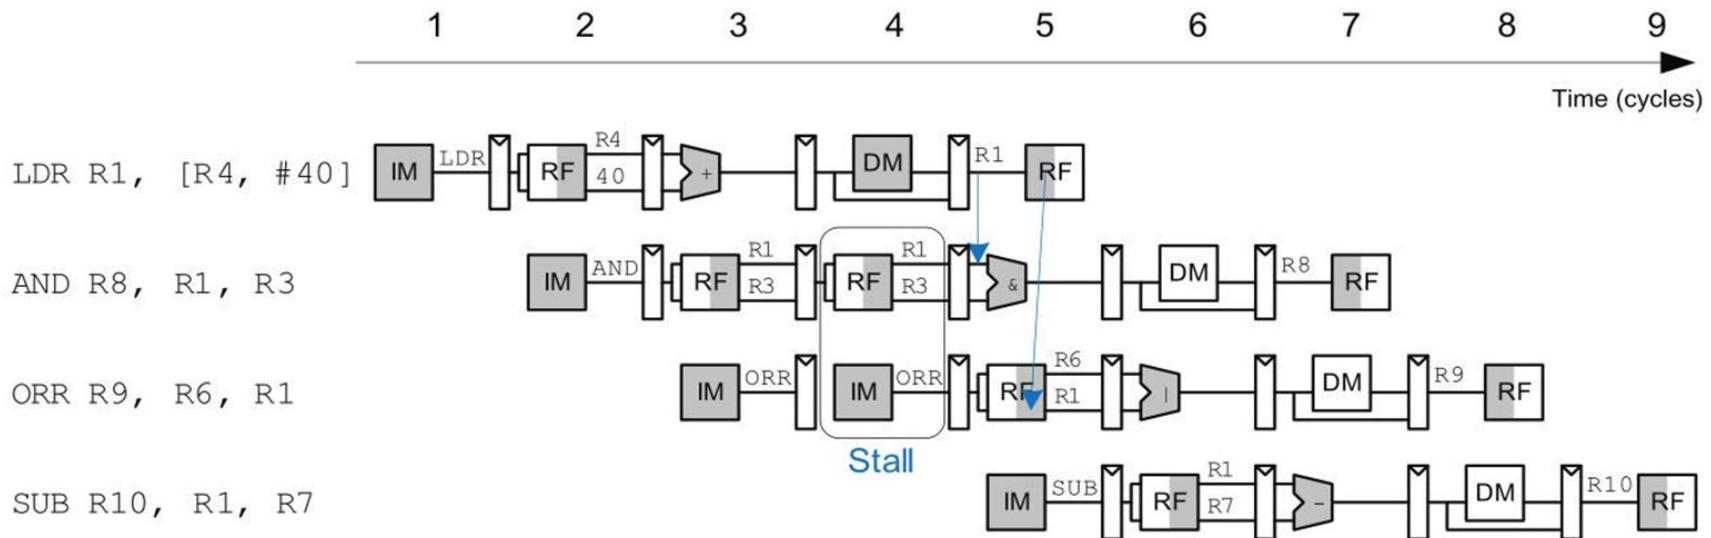


# Pipeline Stalls



- and enters the Decode stage in cycle 3 and stalls there through cycle 4.
- orr must remain in the Fetch stage during both cycles as well, because the Decode stage is full.

# Unused Stages



Execute stage is unused in cycle 4, the Memory stage is unused in cycle 5 and the Writeback stage is unused in cycle 6.

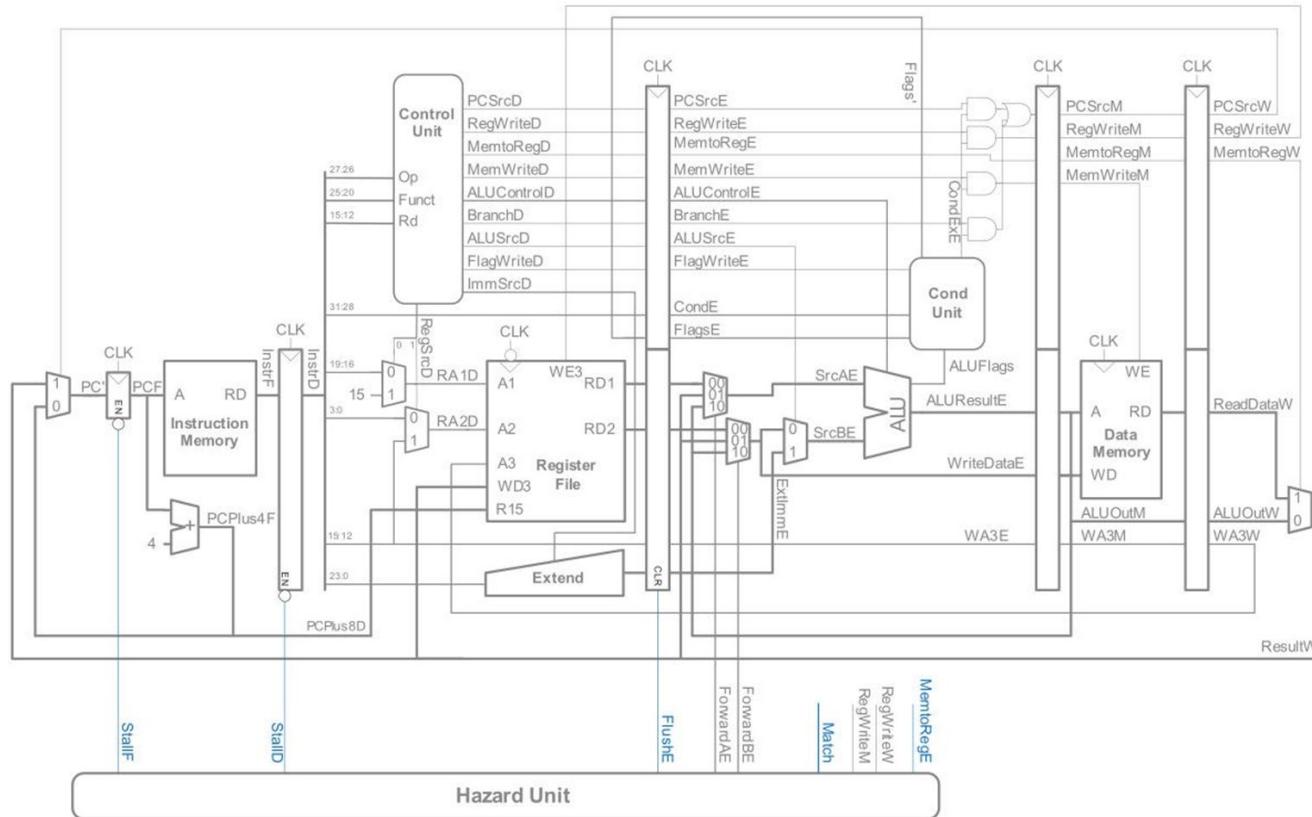
# Bubbles

- This unused stage propagating through the pipeline is called a **bubble**.
- It behaves like a NOP instruction.
- The bubble is introduced by zeroing out the Execute stage control signals during a Decode stall
  - so that the bubble performs no action and changes no architectural state.

# Stalls degrade performance

- Stalling a stage is performed by disabling the pipeline register, so that the contents do not change.
- When a stage is stalled, all previous stages must also be stalled, so that no subsequent instructions are lost.
- The pipeline register directly after the stalled stage must be cleared (flushed) to prevent bogus information from propagating forward.
- Stalls degrade performance, so they should be used only when necessary.

# Pipeline with stalls for ldr



# ldr and Hazard Unit

- The Hazard Unit examines the instruction in the Execute stage.
- If it is an ldr and its destination register (WA3E) matches either source operand of the instruction in the Decode stage (RA1D or RA2D),
  - then that instruction must be stalled in the Decode stage until the source operand is ready.

# Stalls

Stalls are supported by adding enable inputs (EN) to the Fetch and Decode pipeline registers and a synchronous reset/clear (CLR) input to the Execute pipeline register.



# Stall Computation Logic

The logic to compute the stall and flush:

```
Match_12D_E = (RA1D == WA3E) || (RA2D == WA3E)  
LDRstall = Match_12D_E && MemtoRegE  
StallF = StallD = FlushE = LDRstall
```

# Control Hazards

The consequences of B

# Branching instructions

The `b` instruction presents a **control** hazard:

- The pipelined processor does not know what instruction to fetch next
- because the branch decision has not been made by the time the next instruction is fetched.
- Writes to `r15` (PC) present a similar control hazard. (eg, `mov pc,` `lr` or `pop {pc}`)

# B instruction

- One mechanism for dealing with the control hazard is to stall the pipeline until the branch decision is made (i.e., **PCSrcW** is computed).
- Because the decision is made in the Writeback stage, the pipeline would have to be stalled for four cycles at every branch.
- This degrades the system performance if it occurs often.

# PC changes in different phases

The program counter can be changed in different phases

For example:

- bx lr //what phase?
- ldr pc, [sp, #-4] ! //what phase?

As you can see, these PC updates become apparent in different phases of the pipeline

# Branch Prediction

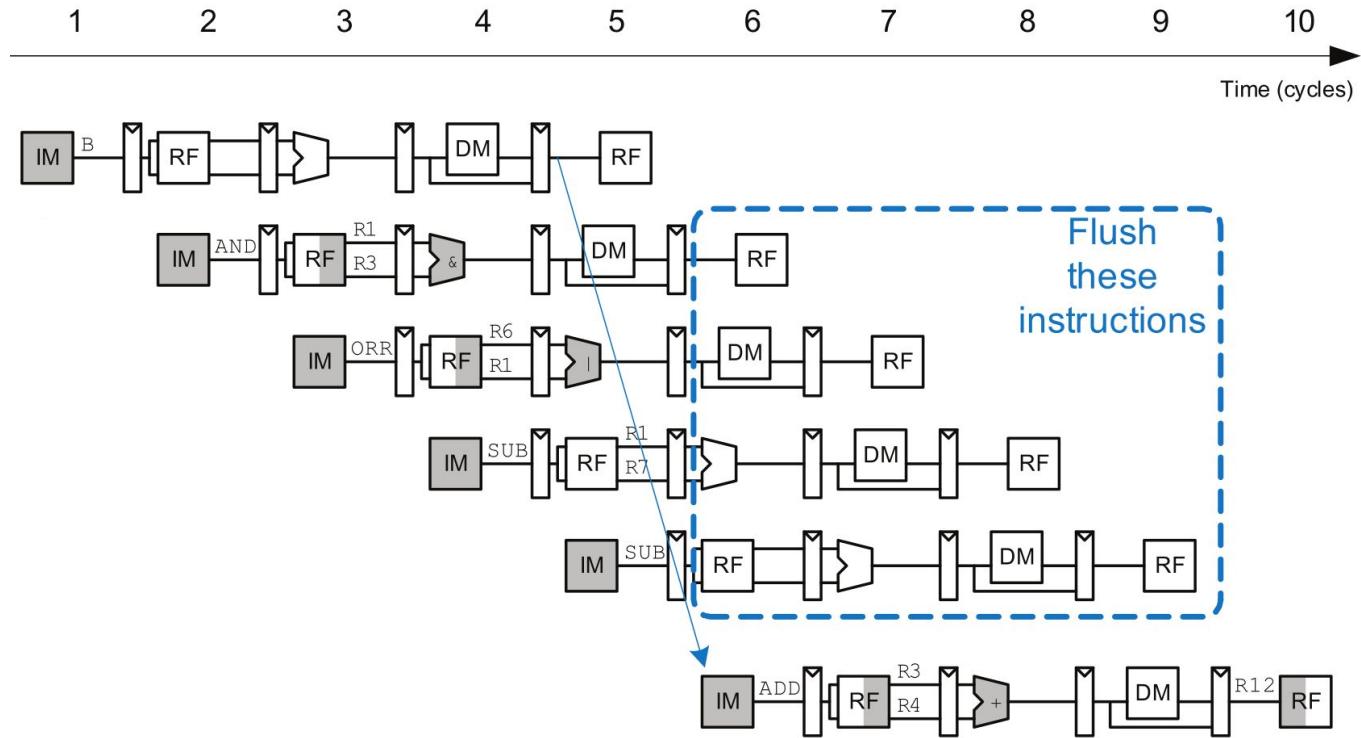
- An alternative is to predict whether the branch will be taken
  - and begin executing instructions based on the prediction.
- Once the branch decision is available, the processor can throw out the instructions
  - if the prediction was wrong.

# Misprediction

- In the pipeline presented so far, the processor predicts that branches are not taken and simply continues executing the program in order until **PCSrcW** is asserted to select the next PC from **ResultW** instead.
- If the branch should have been taken, then the four instructions following the branch must be flushed (discarded) by clearing the pipeline registers for those instructions.
- These wasted instruction cycles are called the branch misprediction penalty.

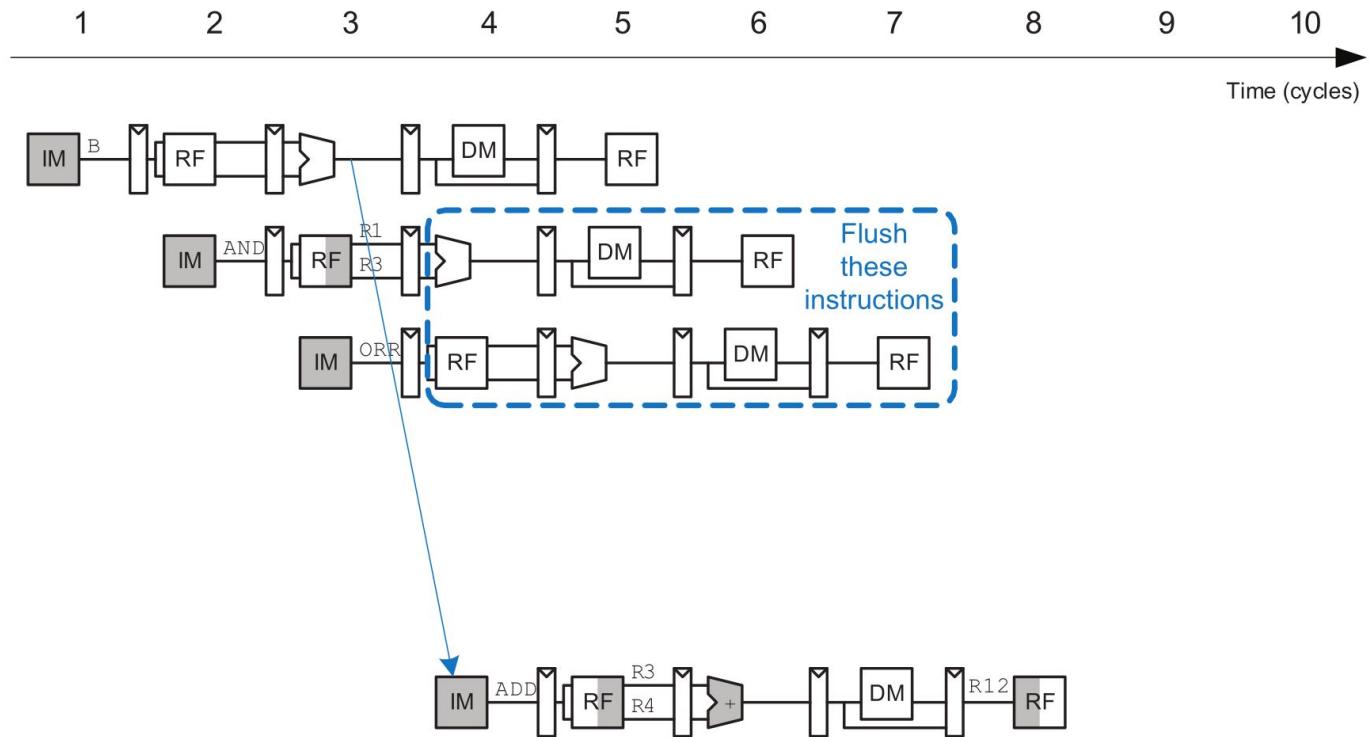
# Flushing on a b

```
b end  
and r8, r1, r3  
orr r9, r6, r1  
sub r10, r1, r7  
sub r11, r1, r8  
...  
end:  
    add r12, r3, r4
```

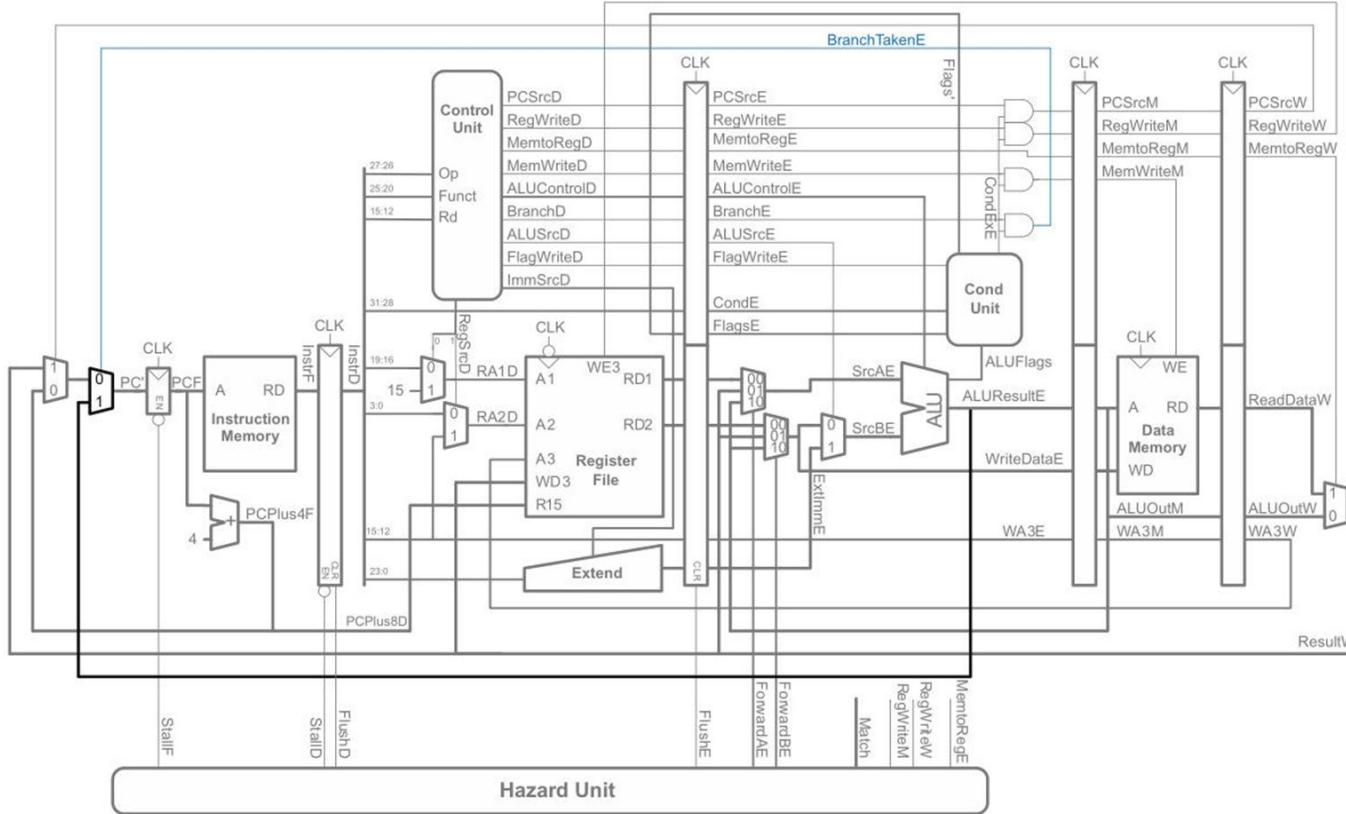


# Execute Stage Decision

```
b end  
and r8, r1, r3  
orr r9, r6, r1  
sub r10, r1, r7  
sub r11, r1, r8  
...  
end:  
    add r12, r3, r4
```



# Branch Control Hazard added



# Branch Multiplexer

- A branch multiplexer is added before the PC register to select the branch destination from **ALUResultE** .
- The **BranchTakenE** signal controlling this multiplexer is asserted on branches whose condition is satisfied.
- **PCSrcW** is now only asserted for writes to the PC, which still occur in the Writeback stage.

# Branch Multiplexer

- When a branch is taken, the subsequent two instructions must be flushed from the pipeline registers of the Decode and Execute stages.
- When a write to the PC is in the pipeline, the pipeline should be stalled until the write completes.
- This is done by stalling the Fetch stage.
- Stalling one stage also requires flushing the next to prevent the instruction from being executed repeatedly.

# Fetch stalled / Decode flushed

- When the PC write reaches the Writeback stage (**PCSrcW** asserted), **StallF** is released to allow the write to occur,
- but **FlushD** is still asserted so that the undesired instruction in the Fetch stage does not advance.

```
PCWrPendingF = PCSrcD || PCSrcE || PCSrcM;
```

```
StallD = LDRstall;
```

```
StallF = LDRstall || PCWrPendingF;
```

```
FlushE = LDRstall || BranchTakenE;
```

```
FlushD = PCWrPendingF || PCSrcW || BranchTakenE;
```

# Reducing the Penalty

- Branches are very common,
- Even a two-cycle misprediction penalty still impacts performance.
- With a bit more work, the penalty could be reduced to one cycle for many branches.
- The destination address must be computed in the Decode stage as
- $\text{PCBranchD} = \text{PCPlus8D} + \text{ExtImmD}$

# Reducing the Penalty

- **BranchTakenD** must also be computed in the Decode stage based on **ALUFlagsE** generated by the previous instruction.
- This might increase the cycle time of the processor if these flags arrive late.

# Pipeline with Full Hazard

