

# CSCI 6461

# Computer Systems

# Architecture

Functions and the Stack

# Labels as functions in ARM32

Call and return

# C function invocation example

3

```
// this is the "callee"  
int find_max(int a, int b, int c) {  
    int max = a;  
    if (b > max) max = b;  
    if (c > max) max = c;  
    return max;  
}  
  
// this is the "caller"  
int main() {  
    int max = find_max(3,44,11);  
    printf("The maximum is %d\n", max);  
}
```

We are all familiar with how function invocations, parameter passing, local variables and return parameters work in C

Let's define the **caller** (`main`) and the **callee** (`find_max`)

# Sequence of Events in Call/Return

4

Call(**caller**):

- (1) prepares function arguments
- (2) remember where to restart
- (3) branch to function
- (8) use result

Inside function (**callee**):

- (4) use method parameters
- (5) perform body of method
- (6) set result
- (7) branch back to restart point

(1) + (4) : where to put method arguments/parameters ?

(2) + (7) : where to put restart point ?

(5) : how to use registers safely ?

(6) + (8) : where to put result ?

# ARM "function" ?

5

- In ARM assembly there is no such thing as a function declaration.
- We have labels that we branch to, and return points we branch back to
- We will soon see how the caller sends parameters to the callee (label)
- And how the callee label returns results (in non-void cases)

# Caller and callee need to agree

6

Parameters: r0 → a, r1 → b, r2 → c

Result → r0

Restart point: r14 (lr - Link Register)

adr lr, restart // (2)

b find\_max // (3)

restart:

    str r0, [r5] // (8)

method ends with:

    mov pc, lr // (7) we will look at this later

# Replace adr and b and mov

7

We can replace the **adr** and **b** instructions with one called **bl** (Branch with Link) which branches to an address stored in a register:

```
bl find_max    //use bl not adr
bx lr          // (7) don't use b
str r0, [r5]   // (8)
```

# Simplified Example

8

```
mov r0, #3           // (1)
mov r1, #44          // (1)
mov r2, #11          // (1)
bl find_max         // (2,3)
str r0, [r5]         // (8)
svc #2
find_max:
. . .
. . .
. . .
bx lr               // (7)
```

Call(**caller**):

- (1) process method arguments
- (2) remember where to restart
- (3) branch to method
- (8) use result

Inside method(**callee**):

- (4) use method parameters
- (5) perform body of method
- (6) set result
- (7) branch back to restart point

```
// this is the "callee"
int find_max(int a, int b, int c) {
    int max = a;
    if (b > max) max = b;
    if (c > max) max = c;
    return max;
}
// this is the "caller"
int main() {
    int max = find_max(3,44,11);
    printf("The maximum is %d\n", max);
}
```

# How to use registers safely

9

- Save any other registers used by method (e.g. just after start) and restore at end “**callee saved**”

```
find_max:  
    str r4, temp //save r4  
    ...  
    //method body can use r4  
    ...  
    ldr r4, temp // restore  
    bx lr  
.data:  
temp .word 0
```

# What if one label invokes another?

10

```
find_max:  
    str lr, temp1  
    bl label_abc  
    // ...  
    ldr pc, temp1  
  
label_abc:  
    str lr, temp2  
    // ...  
    ldr pc, temp2  
.data  
    temp1: .word 0  
    temp2: .word 0
```

The syntax of the `str/ldr` instructions are a syntactical simplification (as mentioned previously)

# Does this scale up?

11

## Recursion

- a method can call any method, including itself
- → maybe one “temp” per register won’t be enough

## Efficiency

- space occupied by all those “temp”s mostly unused at any given moment during run-time
- We need another way to save and restore registers

# Nesting Calls

12

Method calls and returns are nested (like brackets)

- so register saves and restores are also nested
- so what we need is something that works the same way
- Also notice the order in which we save and restore in nested calls

- str r0 ...
- str r1 ...
- ldr r1 ...
- ldr r0 ...

# Introducing the stack

full descending stack

# The Stack

14

A stack is like a very heavy pile of books:

- you can take anything off the top of the pile ("pop")
- you can add anything to the top of the pile ("push")
- you can't move anything into/out of the pile except at the top!

This is simple, and exactly what we need for method calls: **LIFO = last-in, first-out**

# Full / Descending

15

- **Full**: The stack pointer points to the last item pushed onto the stack (not an empty slot).
- **Descending** : The stack grows downward in memory (toward lower addresses).
- Stack and heap grow towards each other
  - This arrangement was devised to simplify memory management
  - It allows room for each of the resources to expand / contract over process running time

# Using a Stack on ARM

16

On most computers:

- stacks start at a large address,
- which decreases as the stack grows

ARM has:

- stack-pointer register (sp , r13)
- 4 different ways of using a stack!


We will use commonest version:

- sp = address of top word on stack (stack-top)
- decrease sp to push, increase to pop

# Stack Indexing Modes

17

- **Push**
  - `str r0, [sp, #-4]!`  
(pre-indexed)
- **Pop**
  - `ldr r0, [sp], #4`  
(post-indexed)



# Preserved (Callee-Saved) Registers

18

- Preserved registers (also called callee-saved registers) are those that a label (callee) must preserve for the caller,
  - the function must restore their original values before returning
- r4–r8, r10, r11: General-purpose registers that a function must save and restore if it modifies them.
- sp (r13): The stack pointer must be preserved in the sense that it must be restored to its original value before returning, ensuring the stack is properly aligned and balanced.

# Preserving r0, r1, r2, r3 ?

19

- If the **caller** needs to preserve r0–r3 it must do so before invoking the label
- The **callee** is not under any obligation to save r0–r3
- **Caller**-saved registers before a label call should use the stack too, eg:

```
str r0, [ sp, #-4 ]! //save r0
bl find_max
ldr r0, [ sp ], #4    //restore r0
```

# Passing Parameters to a label

20

- As `r0, r1, r2, r3` are not preserved registers, the caller should not expect them to be preserved
- They should be used to pass parameters to the function

```
mov r0, #3
mov r1, #44
mov r2, #11
bl find_max
```

```
find_max:
    // can access / over-write r0,..,r3
```

# What if one method calls another? (again)

21

```
find_max:  
    str lr, [sp,#-4]! //push"  
    bl label_abc  
    //some code  
    ldr lr, [sp],#4    //pop"  
    bx lr  
  
label_abc:  
    str lr, [sp,#-4]! //push"  
    bl label_xyz  
    //some code  
    ldr lr, [sp],#4    //pop"  
    bx lr  
  
label_xyz:  
    str lr, [sp,#-4]! //push"  
    // some code  
    ldr lr, [sp],#4    //pop"  
    bx lr
```

You must save the link register on the callee stack (frame) when the label begins

# Brining it together

22

```
mov r0, #3 // prepare param 1
mov r1, #44 // prepare param 2
mov r2, #11 // prepare param 3
bl find_max // invoke using bl => saves the address of svc 2 into lr
svc #2
find_max:
    str lr, [ sp,#-4 ]! // preserve the link registrer
    str r4, [ sp,#-4 ]! // preserve r4 if we plan to use it
    mov r0, #44          // prepare algorithm return value (cheat!)
    ldr r4, [ sp ], #4   // restore r4
    ldr lr, [ sp ], #4   // restore lr
    bx lr                // return
```

# `bx lr` versus `mov pc, lr`

23

- You might sometimes see `mov pc, lr` used interchangeably or in place of `bx lr`
- They are not the same thing
- `bx` is a branching instruction, so *it can be optimized via branch prediction technology within the microarchitecture.*
- Not so with `mov pc, lr` ... so best avoid this, although in practice the performance may be the same.

# Recursion

A function that calls itself

# Recursion

When a function calls itself this pattern is called **recursion**

- We sometimes use recursive implementations of methods
  - because they can lead to compact, elegant code
  - code that may be easier to understand than a corresponding implementation that does not use recursion
- Recursion tends to appear more from the mathematical / analysis viewpoint rather than in commercial software development
- Anything that can be done using recursion, can be done without using recursion
  - By redesigning the algorithm

# Stack Overflow Example

```
// this is a terrible idea
int find_max(int a, int b, int c) {
    int max = a;
    if (b > max) max = b;
    if (c > max) max = c;
    return find_max(a,b,c);
}
```

Segmentation fault (core dumped) ./a.out

```
real-time non-blocking time (microseconds, -R) 200000
core file size (blocks, -c) unlimited
data seg size (kbytes, -d) unlimited
scheduling priority (-e) 0
file size (blocks, -f) unlimited
pending signals (-i) 124653
max locked memory (kbytes, -l) 8192
max memory size (kbytes, -m) unlimited
open files (-n) 1024
pipe size (512 bytes, -p) 8
POSIX message queues (bytes, -q) 819200
real-time priority (-r) 0
stack size (kbytes, -s) 8192
cpu time (seconds, -t) unlimited
max user processes (-u) 124653
virtual memory (kbytes, -v) unlimited
file locks (-x) unlimited
```

# Risks of recursion

- Recursion does not add any computational efficiency or performance advantage to your software
- In fact, it adds **risk** which can destabilise your code
  - As seen above, the runtime does not warn us about the dangers in the code
- Recursive calls have **hard limits** on the stack depth - which is the main constraint that limits its usefulness

# Stack Variables

Making and reclaiming space

# From the last time

29

- Why do we use a stack rather than fixed memory locations?
- A stack is implemented using `sp` and pre-/post-indexing instructions.
  - You should be able to give the code for the **push** and **pop** operations
  - and explain the value in `sp` after each operation.

# Problems

30

Last lecture, we said:

e.g.  $r0 = 1\text{st parameter}$ ,  $r1 = 2\text{nd}$ ,  $r2 = 3\text{rd}$  etc.

- What if we have a lot of parameters?
  - Use the stack if more than four parameters are required (since  $r0-r3$  can only hold up to four 32-bit values).
- Alternative: push arguments onto stack during call
- What if we want to declare some stack local variables?

# Local variables

Local to the stack

# Variable Lifetime

32

Most variables declared inside a function have the same lifetime as the function itself:

- when the function is called, its variables are created
- when the function returns, its variables are destroyed

(there are exceptions, but we will ignore them for now)

# We could use the stack for this

33

The most space-efficient way to do this is:

- at the method start, “**push**” extra space for its variables
- at the method end, “**pop**” the extra space

```
sub sp, sp, #bytes of variables ; create space
```

...

```
add sp, sp, #bytes of variables ; destroy space
```

# Making space for local variables

34

```
// lets make space for
// int max = a;

find_max:
    // after saving lr and r4
    str r0, [ sp, #-4 ]!
```

Of course, we need to restore the stack pointer to the right place at the end of `find_max`

# Create space for 4 local variables

35

We can create room for any number of uninitialized variables by just moving the stack pointer down, for example, four local variable int's:

```
find_max:  
    // save lr and r4  
    sub sp, sp, #16  
    // do the algorithm  
    // restore the 4 ints  
    add sp, sp, #16
```

# Reset the sp and return

36

- Ultimately, at the end of every label, the sp must be reset to the value it had before the label began
- This can be done lots of ways:
  - Pairwise, str ... [sp, #-4], str ... [sp, #-4], str ... [sp, #-4], followed later by ldr ... [sp], #4, ldr ... [sp], #4, ldr ... [sp], #4
  - Or, a bulk reclaim at the end: add sp, sp, #12
  - Just so long as the sp is moved down and up the same number of bytes.

# Bulk Register Load and Store

37

ARM has “Store/Load Multiple” instructions to move several register contents to/from memory with a single instruction

- May apply to any subset of 1-16 registers
- More efficient than using many instructions
- Limited in addressing options: only a single base register

# Store Multiple Full Descending

38

```
// push several registers  
stmfd sp!, { r0, r1, r3-r6 }
```

Store Multiple, Full Descending: STMFD (Store Multiple, Full Descending): stores multiple registers to the stack in a full descending stack model, where `sp` points to the last item pushed and the stack grows downward (toward lower addresses).

# Load Multiple Full Descending

39


```
// pop several registers  
ldmfd sp!, { r0, r1, r3-r6 }
```

LDMFD (Load Multiple, Full Descending) loads multiple registers from the stack in a full descending stack model, where the stack pointer (sp) points to the last item pushed and the stack grows downward.

# Example Stack Entry

40


stmfd r13!, {r4-r7} – pushes r4, r5, r6, r7 onto the stack.



# Example Stack Exit

41

ldmfd r13!, {r4-r7} – pops r4, r5, r6, r7 from the stack.



# Uses

42

Two main uses:

- Saving/restoring registers to create working space
- Copying blocks of memory around

(We are only interested in the first of these.)

**Note**, The `sp` and `pc` can be in the list in ARM instructions, but not in Thumb instructions.

# Saving Registers

43

The assembler has instructions which save typing:

```
push { r0, r1, r3-r6 } // push several registers  
pop { r0, r1, r3-r6 } // pop several registers
```

Which builds a stack in the 'traditional' way.

- The listed register contents are moved (6 in this case)
- The memory used is consecutive; the lowest numbered register always corresponds to the lowest address. (i.e. the top of the stack)

# Save and Restore

44

- If an expression is really complicated, we may need to save/restore registers to evaluate it.
- Can save registers at start of method – “callee saved” (as in previous lecture) or before call – “caller saved”

## Normal Pattern :

- Caller saves when "sending" parameters
- Callee saves when preparing the return branch point `lr` register

# ARM Procedure Call Standard (APCS)

45

- A standard on how to use registers in real programs
- r0–r3: parameter/result passing (extra arguments are stacked) anyone can use, but not saved across call (caller saved)
- r4–r8, r10, r11: temporaries/locals (callee saved)
- r9, r12: temporaries (caller saved)
- sp, lr, pc: special purpose registers

# Stack Frame Example

46

memory address	content	description
0x100	(unused)	initial sp, not used after push
0x9c	parameter	5th parameter (pushed by caller)
0x98	lr	saved return address
0x94	r5	callee-saved register
0x90	r4	callee-saved register
0x8c	local variable 1	function-local data
0x88	local variable 2	function-local data
0x88	[sp points here]	final sp after prologue